

# Chanadana-MiniProject-SLC

March 8, 2020

## 0.1 K-Nearest-Neighbors

KNN falls in the supervised learning family of algorithms. Informally, this means that we are given a labelled dataset consisting of training observations  $(x,y)$  and would like to capture the relationship between  $x$  and  $y$ . More formally, our goal is to learn a function  $h:X \rightarrow Y$  so that given an unseen observation  $x$ ,  $h(x)$  can confidently predict the corresponding output  $y$ .

In this module we will explore the inner workings of KNN, choosing the optimal  $K$  values and using KNN from scikit-learn.

## 0.2 Problem statement

### 0.2.1 Dataset

The data set we'll be using is the Iris Flower Dataset which was first introduced in 1936 by the famous statistician Ronald Fisher and consists of 50 observations from each of three species of Iris (Iris setosa, Iris virginica and Iris versicolor). Four features were measured from each sample: the length and the width of the sepals and petals.

**Source:** <https://archive.ics.uci.edu/ml/datasets/Iris>

**Train the KNN algorithm to be able to distinguish the species from one another given the measurements of the 4 features.**

## 0.3 Question 1

Read the iris.csv file

```
[156]: #Data setup
import pandas as pd

df = pd.read_csv('iris.csv', skiprows=0)
df.sample(10)
```

```
[156]:      Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  \
102  103             7.1             3.0             5.9             2.1
38   39             4.4             3.0             1.3             0.2
93   94             5.0             2.3             3.3             1.0
80   81             5.5             2.4             3.8             1.1
14   15             5.8             4.0             1.2             0.2
30   31             4.8             3.1             1.6             0.2
113  114             5.7             2.5             5.0             2.0
```

139	140	6.9	3.1	5.4	2.1
65	66	6.7	3.1	4.4	1.4
138	139	6.0	3.0	4.8	1.8

	Species
102	Iris-virginica
38	Iris-setosa
93	Iris-versicolor
80	Iris-versicolor
14	Iris-setosa
30	Iris-setosa
113	Iris-virginica
139	Iris-virginica
65	Iris-versicolor
138	Iris-virginica

## 0.4 Data Pre-processing

### 0.5 Question 2 - Estimating missing values

*Its not good to remove the records having missing values all the time. We may end up losing some data points. So, we will have to see how to replace those missing values with some estimated values (median)*

```
[109]: from sklearn.preprocessing import Imputer
imputer = Imputer(missing_values='NaN', strategy='median', axis=0)
imputer = imputer.fit(df.iloc[:, :-1])
imputed_data = imputer.transform(df.iloc[:, :-1].values)
df.iloc[:, :-1] = imputed_data

iris = df
```

```
/home/edwin/anaconda3/lib/python3.7/site-
packages/sklearn/utils/deprecation.py:58: DeprecationWarning: Class Imputer is
deprecated; Imputer was deprecated in version 0.20 and will be removed in 0.22.
Import impute.SimpleImputer from sklearn instead.
  warnings.warn(msg, category=DeprecationWarning)
```

### 0.6 Question 3 - Dealing with categorical data

Change all the classes to numerals (0to2).

```
[110]: iris.iloc[:, 5].unique()
```

```
[110]: array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)
```

```
[111]: iris.head()
```

```
[111]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1.0	5.1	3.5	1.4	0.2	Iris-setosa
1	2.0	4.9	3.0	1.4	0.2	Iris-setosa
2	3.0	4.7	3.2	1.3	0.2	Iris-setosa
3	4.0	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	5.0	3.6	1.4	0.2	Iris-setosa

```
[112]: from sklearn.preprocessing import LabelEncoder
class_label_encoder = LabelEncoder()

iris.iloc[:, -1] = class_label_encoder.fit_transform(iris.iloc[:, -1])
```

```
[113]: iris.head()
```

```
[113]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1.0	5.1	3.5	1.4	0.2	0
1	2.0	4.9	3.0	1.4	0.2	0
2	3.0	4.7	3.2	1.3	0.2	0
3	4.0	4.6	3.1	1.5	0.2	0
4	5.0	5.0	3.6	1.4	0.2	0

## 0.7 Question 4

Observe the association of each independent variable with target variable and drop variables from feature set having correlation in range -0.1 to 0.1 with target variable.

```
[12]: iris.corr()
```

```
[12]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	\
Id	1.000000	0.702734	-0.392693	0.872346	
SepalLengthCm	0.702734	1.000000	-0.109369	0.871120	
SepalWidthCm	-0.392693	-0.109369	1.000000	-0.420713	
PetalLengthCm	0.872346	0.871120	-0.420713	1.000000	
PetalWidthCm	0.890676	0.815986	-0.356510	0.962043	
Species	0.942753	0.775061	-0.417318	0.944477	

	PetalWidthCm	Species
Id	0.890676	0.942753
SepalLengthCm	0.815986	0.775061
SepalWidthCm	-0.356510	-0.417318
PetalLengthCm	0.962043	0.944477
PetalWidthCm	1.000000	0.952513
Species	0.952513	1.000000

## 0.8 Question 5

Observe the independent variables variance and drop such variables having no variance or almost zero variance (variance < 0.1). They will be having almost no influence on the classification.

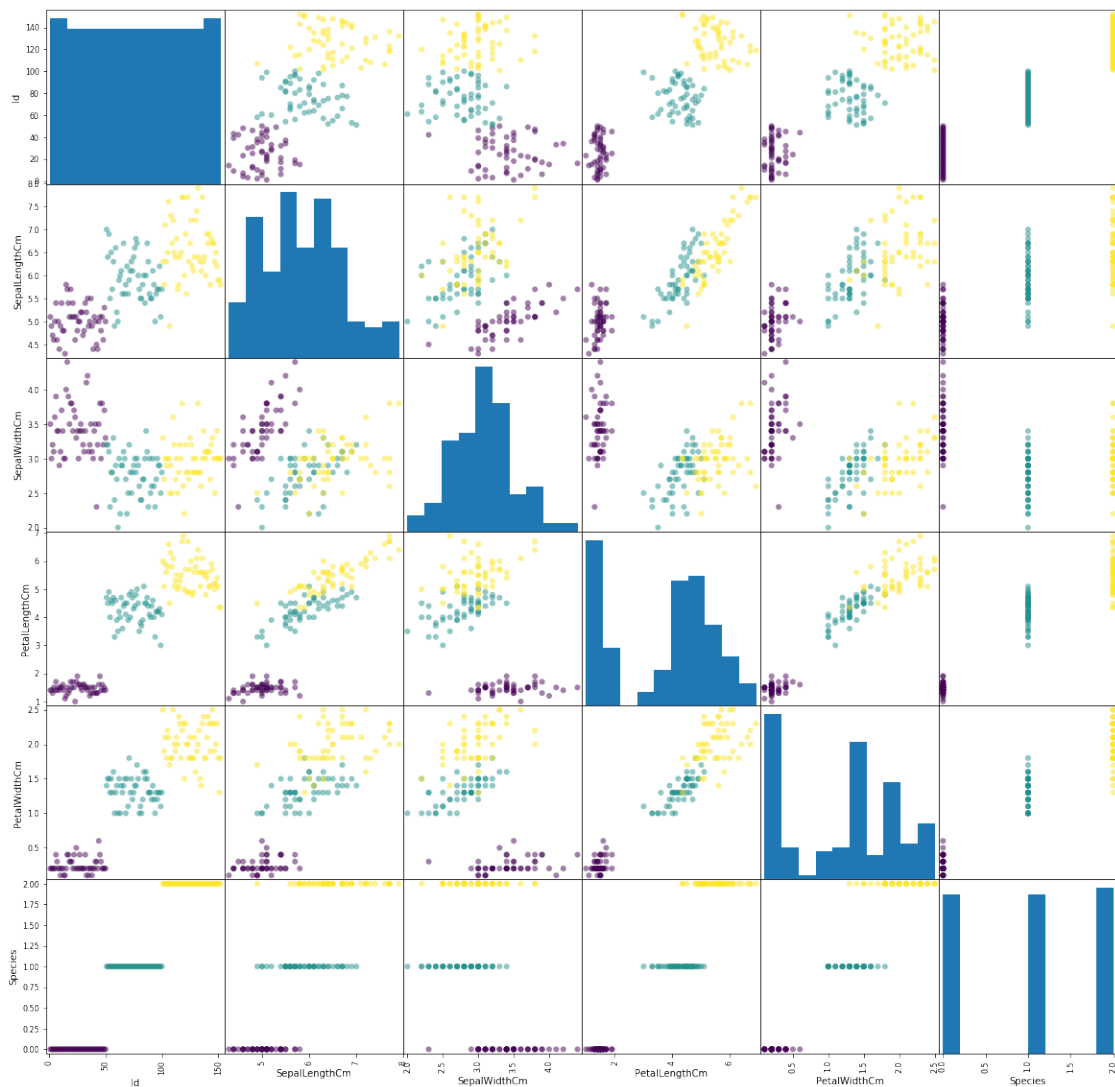
```
[13]: iris.var()
```

```
[13]: Id                1938.000000  
      SepalLengthCm      0.676645  
      SepalWidthCm       0.185552  
      PetalLengthCm      3.076516  
      PetalWidthCm       0.577141  
      Species           0.675322  
      dtype: float64
```

## 0.9 Question 6

*Plot the scatter matrix for all the variables.*

```
[15]: splt = pd.plotting.scatter_matrix(iris, c=iris.iloc[:,-1], figsize=(20, 20),  
    ↪marker='o')
```



## 0.10 Split the dataset into training and test sets

### 0.11 Question 7

*Split the dataset into training and test sets with 80-20 ratio.*

```
[19]: import numpy as np
      from sklearn.model_selection import train_test_split

      # Transform data into features and target
      X = np.array(iris.ix[:, 1:5])
      y = np.array(iris['Species'])

      # split into train and test
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=7)
```

```
[20]: print(X_train.shape)
      print(y_train.shape)
```

```
(121, 4)
(121,)
```

```
[21]: print(X_test.shape)
      print(y_test.shape)
```

```
(31, 4)
(31,)
```

### 0.12 Question 8 - Model

*Build the model and train and test on training and test sets respectively using **scikit-learn**. Print the Accuracy of the model with different values of **k=3,5,9**.*

**Hint:** For accuracy you can check **accuracy\_score()** in scikit-learn

```
[22]: # loading library
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.metrics import accuracy_score

      # instantiate learning model (k = 3)
      knn = KNeighborsClassifier(n_neighbors = 3)

      # fitting the model
      knn.fit(X_train, y_train)

      # predict the response
```

```

y_pred = knn.predict(X_test)

# evaluate accuracy
print(accuracy_score(y_test, y_pred))

# instantiate learning model (k = 5)
knn = KNeighborsClassifier(n_neighbors=5)

# fitting the model
knn.fit(X_train, y_train)

# predict the response
y_pred = knn.predict(X_test)

# evaluate accuracy
print(accuracy_score(y_test, y_pred))

# instantiate learning model (k = 9)
knn = KNeighborsClassifier(n_neighbors=9)

# fitting the model
knn.fit(X_train, y_train)

# predict the response
y_pred = knn.predict(X_test)

# evaluate accuracy
print(accuracy_score(y_test, y_pred))

```

0.9354838709677419

0.967741935483871

0.9032258064516129

### 0.13 Question 9 - Cross Validation

Run the KNN with no of neighbours to be 1,3,5..19 and \*Find the **optimal number of neighbours** from the above list using the Mis classification error

Hint:

Misclassification error (MSE) = 1 - Test accuracy score. Calculated MSE for each model with neighbours = 1,3,5...19 and find the model with lowest MSE

```

[23]: # creating odd list of K for KNN
myList = list(range(1,20))

# subsetting just the odd ones
neighbors = list(filter(lambda x: x % 2 != 0, myList))

```

```
[24]: # empty list that will hold accuracy scores
ac_scores = []

# perform accuracy metrics for values from 1,3,5....19
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    # predict the response
    y_pred = knn.predict(X_test)
    # evaluate accuracy
    scores = accuracy_score(y_test, y_pred)
    ac_scores.append(scores)

# changing to misclassification error
MSE = [1 - x for x in ac_scores]


# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print("The optimal number of neighbors is %d" % optimal_k)
```

The optimal number of neighbors is 5

## 0.14 Question 10

Plot misclassification error vs  $k$  (with  $k$  value on X-axis) using matplotlib.

```
[33]: import matplotlib.pyplot as plt
# plot misclassification error vs k
plt.plot(neighbors, MSE)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()
```



## 1 Naive Bayes

```
[25]: #Load all required library
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
%matplotlib inline
from sklearn import datasets
from sklearn.decomposition import PCA
from sklearn.naive_bayes import GaussianNB, BernoulliNB, MultinomialNB
```

### 1.0.1 Question 1

Import Iris.csv

```
[85]: # Load using input file
iris=pd.read_csv("iris.csv")
iris.head(5)
```

```
[85]:   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
0    1             5.1             3.5             1.4             0.2  Iris-setosa
1    2             4.9             3.0             1.4             0.2  Iris-setosa
2    3             4.7             3.2             1.3             0.2  Iris-setosa
3    4             4.6             3.1             1.5             0.2  Iris-setosa
```



4	5	5.0	3.6	1.4	0.2	Iris-setosa
---	---	-----	-----	-----	-----	-------------

```
[86]: # Check dimension of data
iris.shape
```

```
[86]: (152, 6)
```

```
[87]: #Check shape of data
iris.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 152 entries, 0 to 151
Data columns (total 6 columns):
Id                152 non-null int64
SepalLengthCm     151 non-null float64
SepalWidthCm      150 non-null float64
PetalLengthCm     150 non-null float64
PetalWidthCm      151 non-null float64
Species           152 non-null object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

```
[88]: # check for missing values
```

```
[89]: iris.isna().sum()
```

```
[89]: Id                0
SepalLengthCm        1
SepalWidthCm         2
PetalLengthCm        2
PetalWidthCm         1
Species              0
dtype: int64
```

```
[90]: iris = iris.dropna()
```

```
[91]: iris.isna().sum()
```

```
[91]: Id                0
SepalLengthCm        0
SepalWidthCm         0
PetalLengthCm        0
PetalWidthCm         0
Species              0
dtype: int64
```

### Slice data set for Independent variables and dependent variables

```
X=iris.iloc[:,4].values
y=iris['Species'].values
```

```
[93]: #Check the dataset
print(y)
print(X)
```

10

[ [ 1.	5.1	3.5	1.4]
[ 2.	4.9	3.	1.4]
[ 3.	4.7	3.2	1.3]
[ 4.	4.6	3.1	1.5]
[ 5.	5.	3.6	1.4]
[ 6.	5.4	3.9	1.7]
[ 7.	4.6	3.4	1.4]
[ 8.	5.	3.4	1.5]
[ 9.	4.4	2.9	1.4]
[ 10.	4.9	3.1	1.5]
[ 11.	5.4	3.7	1.5]
[ 12.	4.8	3.4	1.6]
[ 13.	4.8	3.	1.4]
[ 14.	4.3	3.	1.1]
[ 15.	5.8	4.	1.2]
[ 16.	5.7	4.4	1.5]
[ 17.	5.4	3.9	1.3]
[ 18.	5.1	3.5	1.4]
[ 19.	5.7	3.8	1.7]
[ 20.	5.1	3.8	1.5]
[ 21.	5.4	3.4	1.7]
[ 22.	5.1	3.7	1.5]
[ 23.	4.6	3.6	1. ]
[ 24.	5.1	3.3	1.7]
[ 25.	4.8	3.4	1.9]
[ 26.	5.	3.	1.6]
[ 27.	5.	3.4	1.6]
[ 28.	5.2	3.5	1.5]
[ 29.	5.2	3.4	1.4]
[ 30.	4.7	3.2	1.6]
[ 31.	4.8	3.1	1.6]
[ 32.	5.4	3.4	1.5]
[ 33.	5.2	4.1	1.5]
[ 34.	5.5	4.2	1.4]
[ 35.	4.9	3.1	1.5]
[ 36.	5.	3.2	1.2]
[ 37.	5.5	3.5	1.3]
[ 38.	4.9	3.1	1.5]
[ 39.	4.4	3.	1.3]
[ 40.	5.1	3.4	1.5]
[ 41.	5.	3.5	1.3]
[ 42.	4.5	2.3	1.3]
[ 43.	4.4	3.2	1.3]
[ 44.	5.	3.5	1.6]
[ 45.	5.1	3.8	1.9]
[ 46.	4.8	3.	1.4]
[ 47.	5.1	3.8	1.6]
[ 48.	4.6	3.2	1.4]

[ 49.	5.3	3.7	1.5]
[ 50.	5.	3.3	1.4]
[ 51.	7.	3.2	4.7]
[ 52.	6.4	3.2	4.5]
[ 53.	6.9	3.1	4.9]
[ 54.	5.5	2.3	4. ]
[ 55.	6.5	2.8	4.6]
[ 56.	5.7	2.8	4.5]
[ 57.	6.3	3.3	4.7]
[ 58.	4.9	2.4	3.3]
[ 59.	6.6	2.9	4.6]
[ 60.	5.2	2.7	3.9]
[ 61.	5.	2.	3.5]
[ 62.	5.9	3.	4.2]
[ 63.	6.	2.2	4. ]
[ 64.	6.1	2.9	4.7]
[ 65.	5.6	2.9	3.6]
[ 66.	6.7	3.1	4.4]
[ 67.	5.6	3.	4.5]
[ 68.	5.8	2.7	4.1]
[ 69.	6.2	2.2	4.5]
[ 70.	5.6	2.5	3.9]
[ 71.	5.9	3.2	4.8]
[ 72.	6.1	2.8	4. ]
[ 73.	6.3	2.5	4.9]
[ 74.	6.1	2.8	4.7]
[ 75.	6.4	2.9	4.3]
[ 76.	6.6	3.	4.4]
[ 77.	6.8	2.8	4.8]
[ 78.	6.7	3.	5. ]
[ 79.	6.	2.9	4.5]
[ 80.	5.7	2.6	3.5]
[ 81.	5.5	2.4	3.8]
[ 82.	5.5	2.4	3.7]
[ 83.	5.8	2.7	3.9]
[ 84.	6.	2.7	5.1]
[ 85.	5.4	3.	4.5]
[ 86.	6.	3.4	4.5]
[ 87.	6.7	3.1	4.7]
[ 88.	6.3	2.3	4.4]
[ 89.	5.6	3.	4.1]
[ 90.	5.5	2.5	4. ]
[ 91.	5.5	2.6	4.4]
[ 92.	6.1	3.	4.6]
[ 93.	5.8	2.6	4. ]
[ 94.	5.	2.3	3.3]
[ 95.	5.6	2.7	4.2]
[ 96.	5.7	3.	4.2]

[ 97.	5.7	2.9	4.2]
[ 98.	6.2	2.9	4.3]
[ 99.	5.1	2.5	3. ]
[100.	5.7	2.8	4.1]
[101.	6.3	3.3	6. ]
[102.	5.8	2.7	5.1]
[103.	7.1	3.	5.9]
[104.	6.3	2.9	5.6]
[105.	6.5	3.	5.8]
[106.	7.6	3.	6.6]
[107.	4.9	2.5	4.5]
[108.	7.3	2.9	6.3]
[109.	6.7	2.5	5.8]
[110.	7.2	3.6	6.1]
[111.	6.5	3.2	5.1]
[112.	6.4	2.7	5.3]
[113.	6.8	3.	5.5]
[114.	5.7	2.5	5. ]
[115.	5.8	2.8	5.1]
[116.	6.4	3.2	5.3]
[117.	6.5	3.	5.5]
[118.	7.7	3.8	6.7]
[119.	7.7	2.6	6.9]
[120.	6.	2.2	5. ]
[121.	6.9	3.2	5.7]
[122.	5.6	2.8	4.9]
[123.	7.7	2.8	6.7]
[124.	6.3	2.7	4.9]
[125.	6.7	3.3	5.7]
[126.	7.2	3.2	6. ]
[127.	6.2	2.8	4.8]
[128.	6.1	3.	4.9]
[129.	6.4	2.8	5.6]
[130.	7.2	3.	5.8]
[131.	7.4	2.8	6.1]
[132.	7.9	3.8	6.4]
[133.	6.4	2.8	5.6]
[134.	6.3	2.8	5.1]
[135.	6.1	2.6	5.6]
[136.	7.7	3.	6.1]
[137.	6.3	3.4	5.6]
[138.	6.4	3.1	5.5]
[139.	6.	3.	4.8]
[140.	6.9	3.1	5.4]
[141.	6.7	3.1	5.6]
[142.	6.9	3.1	5.1]
[143.	5.8	2.7	5.1]
[144.	6.8	3.2	5.9]

```
[145.    6.7    3.3    5.7]
[146.    6.7    3.    5.2]
[147.    6.3    2.5    5. ]
[148.    6.5    3.    5.2]
[149.    6.2    3.4    5.4]
[150.    5.9    3.    5.1]]
```

### 1.1 Question 3

Find the distribution of target variable (Class)

And, Plot the distribution of target variable using histogram


```
[94]: iris["Species"].value_counts()
```

```
[94]: Iris-setosa      50
      Iris-virginica  50
      Iris-versicolor 50
      Name: Species, dtype: int64
```

#### 1.1.1 Plot the distribution of target variable using histogram

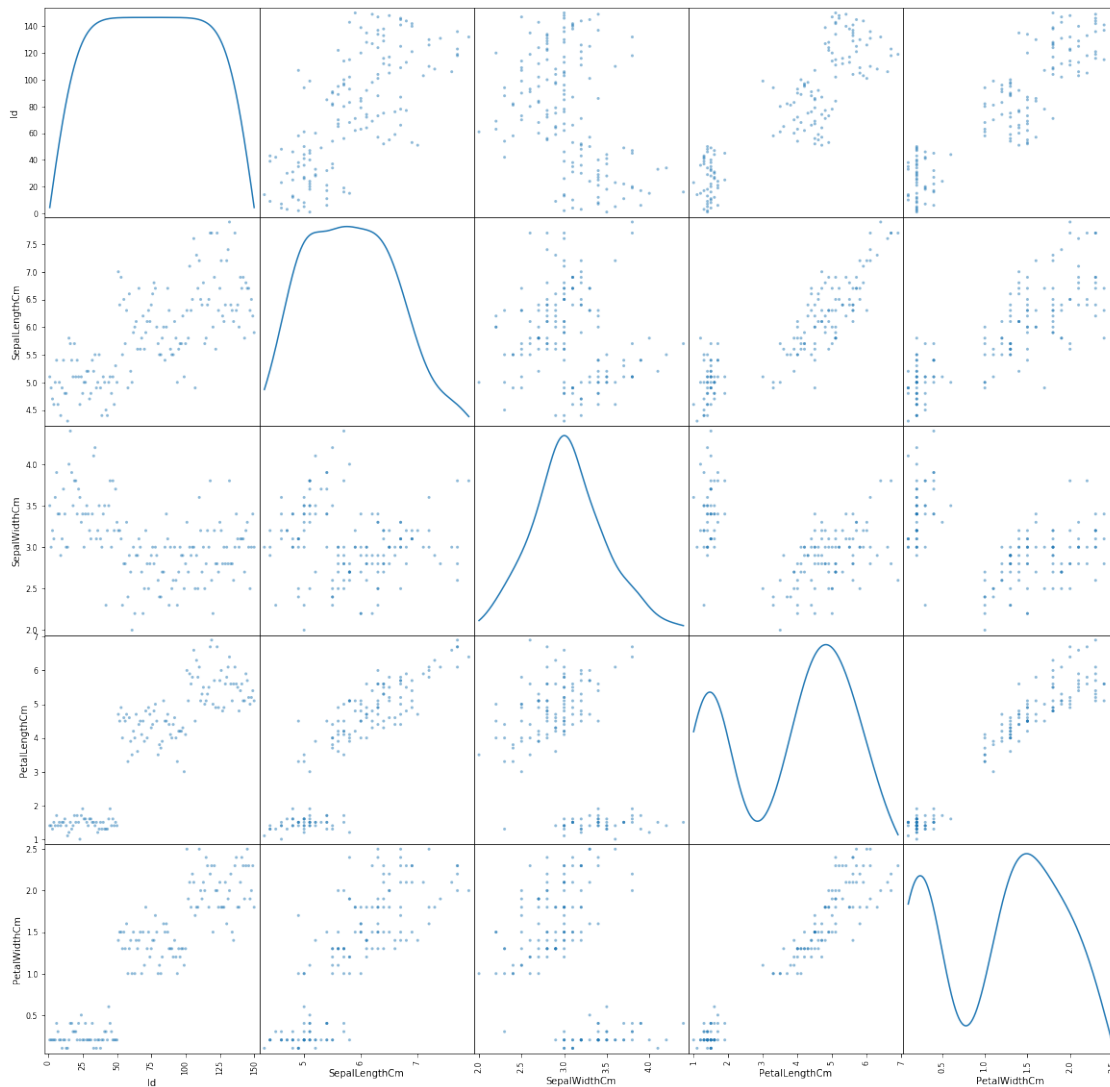
```
[95]: pd.value_counts(iris["Species"]).plot(kind="bar")
```

```
[95]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2d4366c438>
```



### 1.1.2 Plot Scatter Matrix to understand the distribution of variables and give insights from it( 1 Marks)

```
[96]: spd = pd.plotting.scatter_matrix(iris, figsize=(20,20), diagonal="kde")
```



### 1.1.3 Question 3

Find Correlation among all variables and give your insights

```
[97]: corr = iris.corr()
corr
```

*#Please note, it's Require to remove correlated features because they are voted  
 ↳ twice in the model and it can lead to over inflating importance. We will  
 ↳ ignore it here*

```
[97]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	\
Id	1.000000	0.716676	-0.397729	0.882747	
SepalLengthCm	0.716676	1.000000	-0.109369	0.871754	
SepalWidthCm	-0.397729	-0.109369	1.000000	-0.420516	
PetalLengthCm	0.882747	0.871754	-0.420516	1.000000	
PetalWidthCm	0.899759	0.817954	-0.356544	0.962757	

	PetalWidthCm
Id	0.899759
SepalLengthCm	0.817954
SepalWidthCm	-0.356544
PetalLengthCm	0.962757
PetalWidthCm	1.000000

```
[98]: iris
```

```
[98]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	\
0	1	5.1	3.5	1.4	0.2	
1	2	4.9	3.0	1.4	0.2	
2	3	4.7	3.2	1.3	0.2	
3	4	4.6	3.1	1.5	0.2	
4	5	5.0	3.6	1.4	0.2	
5	6	5.4	3.9	1.7	0.4	
6	7	4.6	3.4	1.4	0.3	
7	8	5.0	3.4	1.5	0.2	
8	9	4.4	2.9	1.4	0.2	
9	10	4.9	3.1	1.5	0.1	
10	11	5.4	3.7	1.5	0.2	
11	12	4.8	3.4	1.6	0.2	
12	13	4.8	3.0	1.4	0.1	
13	14	4.3	3.0	1.1	0.1	
14	15	5.8	4.0	1.2	0.2	
15	16	5.7	4.4	1.5	0.4	
16	17	5.4	3.9	1.3	0.4	
17	18	5.1	3.5	1.4	0.3	
18	19	5.7	3.8	1.7	0.3	
19	20	5.1	3.8	1.5	0.3	
20	21	5.4	3.4	1.7	0.2	
21	22	5.1	3.7	1.5	0.4	
22	23	4.6	3.6	1.0	0.2	
23	24	5.1	3.3	1.7	0.5	
24	25	4.8	3.4	1.9	0.2	
25	26	5.0	3.0	1.6	0.2	



26	27	5.0	3.4	1.6	0.4
27	28	5.2	3.5	1.5	0.2
28	29	5.2	3.4	1.4	0.2
29	30	4.7	3.2	1.6	0.2
..	...	...	...	...	...
120	121	6.9	3.2	5.7	2.3
121	122	5.6	2.8	4.9	2.0
122	123	7.7	2.8	6.7	2.0
123	124	6.3	2.7	4.9	1.8
124	125	6.7	3.3	5.7	2.1
125	126	7.2	3.2	6.0	1.8
126	127	6.2	2.8	4.8	1.8
127	128	6.1	3.0	4.9	1.8
128	129	6.4	2.8	5.6	2.1
129	130	7.2	3.0	5.8	1.6
130	131	7.4	2.8	6.1	1.9
131	132	7.9	3.8	6.4	2.0
132	133	6.4	2.8	5.6	2.2
133	134	6.3	2.8	5.1	1.5
134	135	6.1	2.6	5.6	1.4
135	136	7.7	3.0	6.1	2.3
136	137	6.3	3.4	5.6	2.4
137	138	6.4	3.1	5.5	1.8
138	139	6.0	3.0	4.8	1.8
139	140	6.9	3.1	5.4	2.1
140	141	6.7	3.1	5.6	2.4
141	142	6.9	3.1	5.1	2.3
142	143	5.8	2.7	5.1	1.9
143	144	6.8	3.2	5.9	2.3
144	145	6.7	3.3	5.7	2.5
145	146	6.7	3.0	5.2	2.3
146	147	6.3	2.5	5.0	1.9
147	148	6.5	3.0	5.2	2.0
148	149	6.2	3.4	5.4	2.3
149	150	5.9	3.0	5.1	1.8

Species	
0	Iris-setosa
1	Iris-setosa
2	Iris-setosa
3	Iris-setosa
4	Iris-setosa
5	Iris-setosa
6	Iris-setosa
7	Iris-setosa
8	Iris-setosa
9	Iris-setosa

10	Iris-setosa
11	Iris-setosa
12	Iris-setosa
13	Iris-setosa
14	Iris-setosa
15	Iris-setosa
16	Iris-setosa
17	Iris-setosa
18	Iris-setosa
19	Iris-setosa
20	Iris-setosa
21	Iris-setosa
22	Iris-setosa
23	Iris-setosa
24	Iris-setosa
25	Iris-setosa
26	Iris-setosa
27	Iris-setosa
28	Iris-setosa
29	Iris-setosa
..	...
120	Iris-virginica
121	Iris-virginica
122	Iris-virginica
123	Iris-virginica
124	Iris-virginica
125	Iris-virginica
126	Iris-virginica
127	Iris-virginica
128	Iris-virginica
129	Iris-virginica
130	Iris-virginica
131	Iris-virginica
132	Iris-virginica
133	Iris-virginica
134	Iris-virginica
135	Iris-virginica
136	Iris-virginica
137	Iris-virginica
138	Iris-virginica
139	Iris-virginica
140	Iris-virginica
141	Iris-virginica
142	Iris-virginica
143	Iris-virginica
144	Iris-virginica
145	Iris-virginica

```
146 Iris-virginica
147 Iris-virginica
148 Iris-virginica
149 Iris-virginica
```

```
[150 rows x 6 columns]
```

```
[99]: from sklearn.datasets import load_iris
iris = load_iris()


from matplotlib import pyplot as plt

# The indices of the features that we are plotting
x_index = 0
y_index = 1

# this formatter will label the colorbar with the correct target names
formatter = plt.FuncFormatter(lambda i, *args: iris.target_names[int(i)])

plt.figure(figsize=(8, 8))
plt.scatter(iris.data[:, x_index], iris.data[:, y_index], c=iris.target)
plt.colorbar(ticks=[0, 1, 2], format=formatter)
plt.xlabel(iris.feature_names[x_index])
plt.ylabel(iris.feature_names[y_index])

plt.tight_layout()
plt.show()
```



#### 1.1.4 Question 4

Split data in Training and Validation in 80:20

```
[100]: ### SPLITTING INTO TRAINING AND TEST SETS
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.
    ↪ 20, random_state=22)
```

#### 1.1.5 Question 5

Do Feature Scaling

```
[101]: ### NORMALIZATION / FEATURE SCALING
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

### 1.1.6 Question 6

#### Train and Fit NaiveBayes Model

```
[102]: ### WE WILL FIT THE THE CLASSIFIER TO THE TRAINING SET
naiveClassifier=GaussianNB()
naiveClassifier.fit(X_train,y_train)
```

```
[102]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
[103]: y_pred = naiveClassifier.predict(X_test)
```

```
[104]: #Keeping the actual and predicted value side by side
y_compare = np.vstack((y_test,y_pred)).T
#Actual->LEFT
#predicted->RIGHT
#Number of values to be print
y_compare[:20,:]
```

```
[104]: array([[ 'Iris-setosa', 'Iris-setosa'],
 [ 'Iris-virginica', 'Iris-virginica'],
 [ 'Iris-versicolor', 'Iris-versicolor'],
 [ 'Iris-virginica', 'Iris-virginica'],
 [ 'Iris-versicolor', 'Iris-versicolor'],
 [ 'Iris-versicolor', 'Iris-versicolor'],
 [ 'Iris-versicolor', 'Iris-versicolor'],
 [ 'Iris-virginica', 'Iris-virginica'],
 [ 'Iris-versicolor', 'Iris-versicolor'],
 [ 'Iris-setosa', 'Iris-setosa'],
 [ 'Iris-virginica', 'Iris-virginica'],
 [ 'Iris-versicolor', 'Iris-versicolor'],
 [ 'Iris-virginica', 'Iris-virginica'],
 [ 'Iris-virginica', 'Iris-virginica'],
 [ 'Iris-setosa', 'Iris-setosa'],
 [ 'Iris-virginica', 'Iris-virginica'],
 [ 'Iris-versicolor', 'Iris-versicolor'],
 [ 'Iris-versicolor', 'Iris-versicolor'],
 [ 'Iris-virginica', 'Iris-virginica'],
 [ 'Iris-versicolor', 'Iris-versicolor']], dtype=object)
```

### 1.1.7 Question 7

## Print Accuracy and Confusion Matrix and Conclude your findings

```
[157]: # Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```

-----
ValueError                                Traceback (most recent call
last)

<ipython-input-157-4356c19c44e3> in <module>
      1 # Making the Confusion Matrix
      2 from sklearn.metrics import confusion_matrix
----> 3 cm = confusion_matrix(y_test, y_pred)
      4 print(cm)

~/anaconda3/lib/python3.7/site-packages/sklearn/metrics/classification.
py in confusion_matrix(y_true, y_pred, labels, sample_weight)
    251
    252     """
--> 253     y_type, y_true, y_pred = _check_targets(y_true, y_pred)
    254     if y_type not in ("binary", "multiclass"):
    255         raise ValueError("%s is not supported" % y_type)

~/anaconda3/lib/python3.7/site-packages/sklearn/metrics/classification.
py in _check_targets(y_true, y_pred)
     69     y_pred : array or indicator matrix
     70     """
--> 71     check_consistent_length(y_true, y_pred)
     72     type_true = type_of_target(y_true)
     73     type_pred = type_of_target(y_pred)

~/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py in
check_consistent_length(*arrays)
    233     if len(uniques) > 1:
    234         raise ValueError("Found input variables with inconsistent
numbers of"
--> 235                               " samples: %r" % [int(l) for l in lengths])
    236
    237
```

ValueError: Found input variables with inconsistent numbers of samples:  
↪ [192, 30]

```
[106]: #finding accuracy from the confusion matrix.
a = cm.shape
correctPrediction = 0
falsePrediction = 0

for row in range(a[0]):
    for c in range(a[1]):
        if row == c:
            correctPrediction += cm[row,c]
        else:
            falsePrediction += cm[row,c]
print('Correct predictions: ', correctPrediction)
print('False predictions', falsePrediction)
print ('\n\nAccuracy of the Naive Bayes Clasification is: ', correctPrediction/
↪ (cm.sum()))
```

Correct predictions: 30  
False predictions 0

Accuracy of the Naive Bayes Clasification is: 1.0

```
[107]: from sklearn import metrics
print(metrics.classification_report(y_pred, y_test))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	6
Iris-versicolor	1.00	1.00	1.00	10
Iris-virginica	1.00	1.00	1.00	14
micro avg	1.00	1.00	1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

## 2 SVM

```
[158]: #Import library
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
%matplotlib inline
```

This dataset describes the medical records for Pima Indians and whether or not each patient will have an onset of diabetes within

ve years.

Fields description follow:

preg = Number of times pregnant

plas = Plasma glucose concentration a 2 hours in an oral glucose tolerance test

pres = Diastolic blood pressure (mm Hg)

skin = Triceps skin fold thickness (mm)

test = 2-Hour serum insulin (mu U/ml)

mass = Body mass index (weight in kg/(height in m)<sup>2</sup>)

pedi = Diabetes pedigree function

age = Age (years)

class = Class variable (1:tested positive for diabetes, 0: tested negative for diabetes)

### 2.0.1 Question 1

Read the input file 'Diabetes.csv' using Pandas and check it's column names(1 Marks)

```
[159]: diabetes = pd.read_csv('pima-indians-diabetes.csv')
print(diabetes.columns)
```

```
Index(['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class'],
      dtype='object')
```

```
[160]: # Eye ball the imported dataset
diabetes.head()
```

```
[160]:
```

	preg	plas	pres	skin	test	mass	pedi	age	class
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

### 2.0.2 Question 2

Check the dimensions of dataset

```
[161]: print("dimension of diabetes data: {}".format(diabetes.shape))
#The diabetes dataset consists of 768 data points, with 9 features
```

```
dimension of diabetes data: (768, 9)
```



### 2.0.3 Question 3

Check distribution of dependent variable 'class' and plot it

```
[162]: print(diabetes.groupby('class').size())
```

```
class
0     500
1     268
dtype: int64
```


2.0.4 Out of 768 data points, 500 are labeled as 0 and 268 as 1.

2.0.5 Class 0 means No diabetes, outcome 1 means diabetes

```
[163]: import seaborn as sns

sns.countplot(diabetes['class'],label="Count")
```

```
[163]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2d41287e10>
```



```
[164]: diabetes.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
preg      768 non-null int64
```


```
plas      768 non-null int64
pres      768 non-null int64
skin      768 non-null int64
test      768 non-null int64
mass      768 non-null float64
pedi      768 non-null float64
age       768 non-null int64
class     768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

## 2.0.6 Question 4

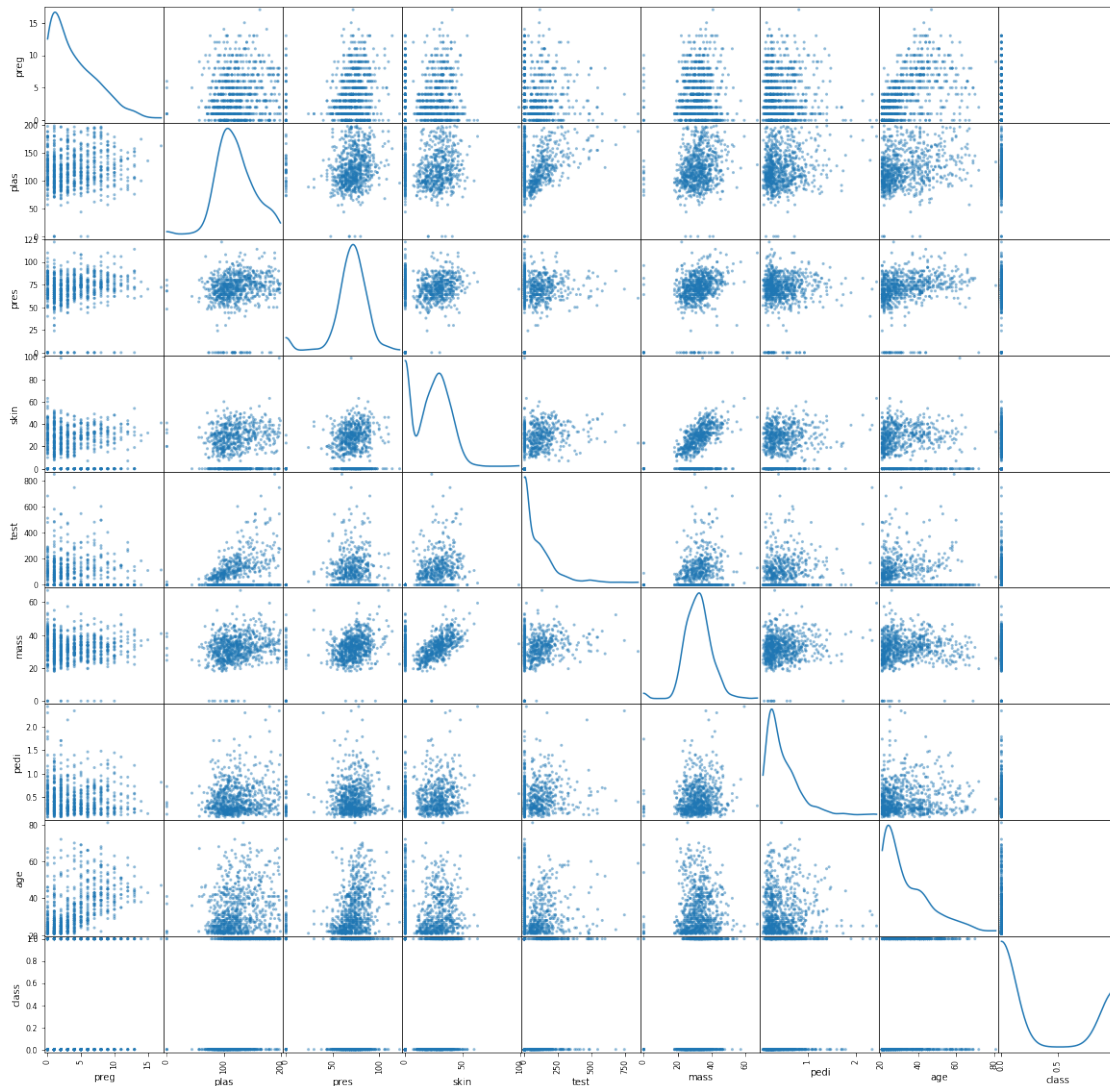
Do correlation analysis and bivariate viualization with Insights

```
[165]: colormap = plt.cm.viridis # Color range to be used in heatmap
plt.figure(figsize=(15,15))
plt.title('Pearson Correlation of attributes', y=1.05, size=19)
sns.heatmap(diabetes.corr(),linewidths=0.1,vmax=1.0,
            square=True, cmap=colormap, linecolor='white', annot=True)
#There is no strong correlation between any two variables.
#There is no strong correlation between any independent variable and class_
↪variable.
```

```
[165]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2d386d70f0>
```



```
[166]: spd = pd.plotting.scatter_matrix(diabetes, figsize=(20,20), diagonal="kde")
```



### 2.0.7 Question 5

Do train and test split with stratify sampling on Outcome variable to maintain the distribution of dependent variable

```
[167]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(diabetes.loc[:, diabetes.
↳columns != 'class'], diabetes['class'], stratify=diabetes['class'],
↳random_state=11)
```

```
[168]: X_train.shape
```

```
[168]: (576, 8)
```

### 2.0.8 Question 6

#### Train Support Vector Machine Model

```
[169]: from sklearn.svm import SVC

svc = SVC()
svc.fit(X_train, y_train)

print("Accuracy on training set: {:.2f}".format(svc.score(X_train, y_train)))
print("Accuracy on test set: {:.2f}".format(svc.score(X_test, y_test)))
```

Accuracy on training set: 1.00

Accuracy on test set: 0.65

```
[170]: #The model overfits substantially with a perfect score on the training set and
      ↪ only 65% accuracy on the test set.

      #SVM requires all the features to be on a similar scale. We will need to
      ↪ rescale our data that all the features are approximately on the same scale
      ↪ and then see the performance
```

### 2.0.9 Question 7

#### Scale the data points using MinMaxScaler

```
[171]: from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.fit_transform(X_test)
```

```
/home/edwin/anaconda3/lib/python3.7/site-
packages/sklearn/preprocessing/data.py:323: DataConversionWarning: Data with
input dtype int64, float64 were all converted to float64 by MinMaxScaler.
```

```
    return self.partial_fit(X, y)
/home/edwin/anaconda3/lib/python3.7/site-
packages/sklearn/preprocessing/data.py:323: DataConversionWarning: Data with
input dtype int64, float64 were all converted to float64 by MinMaxScaler.
    return self.partial_fit(X, y)
```

### 2.0.10 Question 8

#### Fit SVM Model on scaled data and give your observation

```
[172]: svc = SVC()
svc.fit(X_train_scaled, y_train)

print("Accuracy on training set: {:.2f}".format(svc.score(X_train_scaled,
      ↪ y_train)))
```

```
print("Accuracy on test set: {:.2f}".format(svc.score(X_test_scaled, y_test)))
```

Accuracy on training set: 0.76

Accuracy on test set: 0.79

### 2.0.11 Question 9

Try improving the model accuracy using C=1000

```
[173]: svc = SVC(C=1000)
      svc.fit(X_train_scaled, y_train)

      print("Accuracy on training set: {:.3f}".format(
          svc.score(X_train_scaled, y_train)))
      print("Accuracy on test set: {:.3f}".format(svc.score(X_test_scaled, y_test)))
```

Accuracy on training set: 0.812

Accuracy on test set: 0.755